Differentially private vector aggreation in the case of multivariate gaussian data

Tim Sehested Poulsen

December 28, 2022

1 Preliminaries

1.1 Definitions

The dataset is denoted by X, where $X \in \mathbb{R}^{n \times d}$. I have that n denotes the number of entries in the dataset and d is the number of dimensions of the dataset. I will throughout the report refer to a single entry of the dataset as x_i and a single dimension of the dataset as $X^{(j)}$, and therefore $x_i^{(j)}$ denotes the j'th dimension of the i'th entry.

Differential privacy is the heuristic of releasing a database statistic whilst limiting the impact of any one entry. It builds on the intuition that computing a statistics on a private dataset should not reveal any sensitive information about any one individual as long as that individual has little to no effect on the outcome. Differential privacy has multiple slightly different formal definitions, one such is (ε, δ)-Differential Privacy refered to as (ε, δ)-DP which will be introduced later on. A prerequisite for almost all of the different differential privacy definitions relies on the concept of neighbouring dataset.

Definition 1.1 (Neighbouring dataset [5]). *Two dataset $X, X' \in \mathbb{R}^{n \times d}$ are said to be neigbouring if they differ in at most a single entry. Neighbouring dataset are denoted with the relation $X \sim X'$ and defined as followed*

$$X \sim X' \iff |\{i \in \mathbb{N} \mid i \leq n \wedge x_i \neq x_i'\}| \leq 1$$

Definition 1.2 (Sensitivity [6]). *Let $f(X) : \mathbb{R}^{n \times d} \to \mathbb{R}^k$ be a function. The l_p-sensitivity of f is the maximal possible l_p-norm of the difference between the output of f on two neighbouring dataset. We denote the sensitivity as*

$$\Delta_p(f) = \max_{X \sim X'} \|f(X) - f(X')\|_p$$

and then the total l_2-sensitivity is then

Throughout the report I will only be working with l_2-sensitivity and will just denote this as $\Delta(f)$ for ease of notation.

Definition 1.3 ((ε, δ)-Differential Privacy [5]). *A randomized algorithm $\mathcal{M} : \mathbb{R}^{n \times d} \to \mathcal{R}$ is (ε, δ)-differentially private if for all possible subsets of outputs $S \subseteq \mathcal{R}$ and all pairs of neighbouring dataset $X \sim X'$ we have that*

$$\Pr\left[M(X) \in S\right] \le e^{\varepsilon} \cdot \Pr\left[M(X') \in S\right] + \delta$$

Theorem 1 ((ε, δ)-DP under post-processing [6]). *Let $\mathcal{M} : \mathbb{R}^{n \times d} \to \mathcal{R}$ be an (ε, δ)-DP algorithm. Let $f : \mathcal{R} \to \mathcal{R}'$ be an arbitrary mapping, then $f \circ \mathcal{M} : \mathbb{R}^{n \times d} \to \mathcal{R}'$ is (ε, δ)-DP* .

   Proof

   Fix any pair of neighbouring datasets $X \sim X'$ and let $S \subseteq \mathcal{R}'$ be an arbitrary event. We then define $T = \{r \in \mathcal{R} \mid f(r) \in S\}$. We thus have that

$$\begin{aligned}
\Pr\left[f(\mathcal{M}(X)) \in S\right] &= \Pr\left[\mathcal{M}(X) \in T\right] \\
&\le e^{\varepsilon} \cdot \Pr\left[\mathcal{M}(X') \in T\right] + \delta = e^{\varepsilon} \cdot \Pr\left[f(\mathcal{M}(X')) \in S\right] + \delta
\end{aligned}$$

∎

Error Measure   As this report concerns itself exclusively with the sum of entries in a dataset, error will be defined as the expected squared l_2-norm between the true sum and the output of a randomized algorithm. So let $X \in \mathbb{R}^{n \times d}$ be the dataset and $f(X) = \sum_i^n x_i$ be the true sum of all entries. The error of a randomized algorithm $M : \mathbb{R}^{n \times d} \to \mathbb{R}^d$ which estimates $f(X)$ is then

$$\mathrm{Err}(M) := \mathbb{E}\left[\|M(X) - f(X)\|^2\right]$$

extra

> If proof of elliptical ... is included then write about edp is preserved during transformation

1.2   Quadratic forms of random variables

Quadratics of random variables have been well studied [2, 10], specially in the case of multivariate gaussian varaibles [7, 10]. Even more research has been done in evaluating the CDF of these quadratic forms for Gaussian random vectors [4, 8].

Theorem 2 (Expectation of a quadratic random variable [2]). *Let X be a d-dimensional random vector with expected value $\mathbb{E}\left[X\right] = \boldsymbol{\mu_X}$ and covariance matrix $Var\left[X\right] = \boldsymbol{\Sigma_X}$. Let also A be a constant $d \times d$ symmetric matrix, then*

$$\mathbb{E}\left[X^T A X\right] = tr\left(A\boldsymbol{\Sigma_X}\right) + \boldsymbol{\mu}^T A \boldsymbol{\mu}$$

   Proof
Blah blah

$$\begin{aligned}
\mathbb{E}\left[X^T A X\right] &= \mathrm{tr}\left(\mathbb{E}\left[X^T A X\right]\right) = \mathbb{E}\left[\mathrm{tr}\left(X^T A X\right)\right] \\
&= \mathbb{E}\left[\mathrm{tr}\left(A X X^T\right)\right] = \mathrm{tr}\left(A\mathbb{E}\left[X X^T\right]\right) = \mathrm{tr}\left(A\left(Var\left[X\right] + \boldsymbol{\mu}\boldsymbol{\mu}^T\right)\right) \\
&= \mathrm{tr}\left(A\boldsymbol{\Sigma}\right) + \mathrm{tr}\left(A\boldsymbol{\mu}\boldsymbol{\mu^T}\right) = \mathrm{tr}\left(A\boldsymbol{\Sigma}\right) + \boldsymbol{\mu}^T A \boldsymbol{\mu}
\end{aligned}$$

blah

∎

Corollary 1.1. *Let $X \sim \mathcal{N}(\mathbf{0}, \Sigma_{\boldsymbol{X}})$ be a d-dimensional gaussian vector with expected value $\mathbf{0}$, and let σ_j^2 denote the variance of the j'th dimension where $1 \leq j \leq d$. By theorem 2 we have that the expected l_2-norm of such a vector is given by*

$$\mathbb{E}\left[\|X\|^2\right] = tr(\Sigma_{\boldsymbol{X}}) = \sum_{j=1}^{d} \sigma_j^2$$

2 Algorithms

2.1 The Gaussian Mechanism

One of the most foundational algorithms for achieving (ε, δ)-DP is the Gaussian Mechanism [6]. It computes the real value of a statistic, where the l_2-sensitivity is known. That is it produces a (ε, δ)-DP estimate of a function $g : \mathbb{R}^{n \times d} \to \mathbb{R}^d$ where the l_2-sensitivity $\Delta(g)$ is known. It does so by computing the value of $g(X)$ and then adding noise to each dimension drawn from the normal distribution $\mathcal{N}(0, \sigma_{\varepsilon,\delta}^2)$. This can be seen as adding a noise vector η which is then distributed according to the multivariate normal distribution $\mathcal{N}(\mathbf{0}, \sigma_{\varepsilon,\delta}^2 I)$. The algorithm can be seen in Algorithm 1.

Algorithm 1 The Gaussian Mechanism

Input

    $\sigma_{\varepsilon,\delta}$ Standard deviation required to achieve (ε, δ)-DP

    $X \in \mathbb{R}^{n \times d}$ Dataset

Output

    (ε, δ)-DP estimate of $g(X)$

$\eta \leftarrow$ sample from $\mathcal{N}(\overrightarrow{0}, \sigma_{\varepsilon,\delta}^2 I)$

return $g(X) + \eta$

The algorithm quite intuitively prodcues error which is purely given by the norm of the noise added, and the expected error can be calculated to be

$$\mathbb{E}\left[\|(g(X) + \eta) - g(X)\|^2\right] = \mathbb{E}\left[\|\eta\|^2\right]$$

Which by corollary 1.1 is

$$\mathbb{E}\left[\|\eta\|^2\right] = \sum_{i}^{d} \sigma_{\varepsilon,\delta}^2 = d \cdot \sigma_{\varepsilon,\delta}^2 \tag{1}$$

It is apparent that the main difficulty of the mechanism lies in determining a $\sigma_{\varepsilon,\delta}$ which achieves (ε, δ)-DP , and preferably the smallest such one.

    The following theorem was initially proven

Theorem 3. [6] *Let $g : \mathbb{R}^{n \times d} \to \mathbb{R}^d$ be an arbitrary d-dimensional function with l_2-sensitvity $\Delta(g) = \max_{X \sim X'} \|g(X) - g(X')\|$, and let $\varepsilon \in (0, 1)$. The Gaussian Mechanism with $\sigma_{\varepsilon,\delta} = \Delta(g)\sqrt{2\ln(1.25/\delta)}/\varepsilon$ is (ε, δ)-DP .*

    The proof is rather long and is therefore ommitted here. A few years later it was shown in [1] how to compute the minimal $\sigma_{\varepsilon,\delta}$.

Theorem 4. [1] *The Gaussian Mechanism is differentially private if and only if $\sigma_{\varepsilon,\delta} \geq \Delta(g) \cdot \sigma_{opt}$ where $\sigma_{opt} \in \mathbb{R}$ is the smallest value greater than 0, which satisfies*

$$\Phi\left(\frac{1}{2\sigma_{opt}} - \varepsilon\sigma_{opt}\right) - e^{\varepsilon}\Phi\left(-\frac{1}{2\sigma_{opt}} - \varepsilon\sigma_{opt}\right) \leq \delta$$

In [1] it is also shown how to compute this value, and since this is of no importance to this project it will therefore not be covered. For the rest of the report I will only be reffering to $\sigma_{\varepsilon,\delta}$ as the minimal value given by theorem 4 and not that given by theorem 3.

The main downside of the Gaussian Mechanism is that it adds equal noise to all dimensions, regardless of the sensitivity in that dimension. As this report focuses on giving differentially private estimates of sums of vectors, it is natural to ask whether adding equal noise in all dimensions is optimal in this setting. Let us define the function of interest $f(X) = \sum_{i \in [n]} x_i$, for a dataset $X \in \mathbb{R}^{n \times d}$. The sensitivity of this function must be given by the largest possible norm of any vector. Therefore if the largest difference between any two neighbouring datasets in the j'th dimension is given by

$$\Delta_j := \max_{X \sim X'} |X^{(j)} - X'^{(j)}| \tag{2}$$

and we say that X and X' differ in the i'th entry we conclude that

$$\Delta(f) = \max_{X \sim X'} \|f(X) - f(X')\| = \max_{X \sim X'} \|x_i - x_i'\| = \sqrt{\sum_{j \in [d]} \Delta_j^2} = \|\boldsymbol{\Delta}\| \tag{3}$$

This means that by equation 1 the expected error of the Gaussian Mechanism when estimating $f(X)$ is given by

$$\mathbb{E}\left[\|\eta\|^2\right] = d \cdot \sigma_{\varepsilon,\delta}^2 = d \cdot (\Delta(f)\sigma_{opt})^2 = d \cdot \sigma_{opt}^2 \cdot \|\boldsymbol{\Delta}\|^2 \tag{4}$$

A logical next step would be to add noise to each dimensions such that it is proportional to the sensitivity of that dimension. This has been studied in [11] and lays the foundation for this report. Their mechanism, appropriately called the Elliptical Gaussian Mechanism works very similairly to the Gaussian Mechanism as described by algorithm 1, though instead of sampling η_j from $\mathcal{N}(0, \sigma_{\varepsilon,\delta}^2)$, here it is instead drawn from $\mathcal{N}\left(0, \left(\sigma_{opt} \cdot \frac{\Delta_j}{b_j}\right)^2\right)$. The algorithm is described in detail in Algorithm 2

> add lemma/theorem/proof of why this is (ε, δ)-DP (figure out if needed). Prolly is as the method is much the same for gaussian data

The main thing that really differentiates itself from the normal Gaussian Mechanism is the introduction of the scaling vector $\boldsymbol{b}$. This is the scaling of how much weight should be attributed to each dimension when adding noise. It is shown how to determine the optimal values for b_j in [11], and the expected error of the mechanism.

Theorem 5 (Optimality and error of the Elliptical Gaussian Mechanism [11])**.** *The value for b_j which minimized expected l_2 error $\mathbb{E}\left[\|\eta\|^2\right]$ of the Elliptical Gaussian Mechanism is as follows*

$$b_j = \sqrt{\frac{\Delta_j}{\sum_{j \in [d]} \Delta_j}}$$

Algorithm 2 The Elliptical Gaussian Mechanism

Input

σ_{opt} Standard deviation as defined by theorem 4
$X \in \mathbb{R}^{n \times d}$ Dataset
$\boldsymbol{b} \in \mathbb{R}^d$ Scaling vector, where $\|\boldsymbol{b}\| = 1$
$\boldsymbol{\Delta} \in \mathbb{R}^d$ Sensitivies of all dimensions

Output

(ε, δ)-DP estimate of $f(X)$

for $j \in [d]$ **do**
    $\sigma_j \leftarrow \sigma_{opt} \cdot \frac{\Delta_j}{b_j}$
    $\eta_j \leftarrow$ sample from $\mathcal{N}(0, \sigma_j^2)$
end for
return $g(X) + \eta$

---

which leads the error to be

$$\mathbb{E}\left[\|\eta\|^2\right] = \sigma_{opt} \cdot \|\boldsymbol{\Delta}\|_1^2 \tag{5}$$

where $\|\cdot\|_1$ *is the* l_1 *norm.*

> write the proof or skip it, it is going to be very similair to that which I will make

     Comparing the expected error between Algorithm 1 and Algorithm 2 we have the following ratio

$$\frac{d \cdot \sigma_{opt}^2 \cdot \|\boldsymbol{\Delta}\|^2}{\sigma_{opt} \cdot \|\boldsymbol{\Delta}\|_1^2} = \left(\frac{\sqrt{d}\|\boldsymbol{\Delta}\|}{\|\boldsymbol{\Delta}\|_1}\right)^2$$

in which it can be seen that they are equal when all entries of $\boldsymbol{\Delta}$ are the same. Otherwise the error for the Elliptical Gaussian Mechanism is lower when $\boldsymbol{\Delta}$ is skewed. As argued in [11] Algorithm 2 improves Algorithm 1 by a factor in $[1, d)$.

3   Problem setup

As previously mentioned the problem investigated here consists of realeasing the sum of vectors in a dataset under differential privacy. More formally we whish to release the value of $f(X) = \sum_{i=1}^n x_i$ under (ε, δ)-DP .
The common factor for achieving (ε, δ)-DP in both the Gaussian Mechanism and the Elliptical Gaussian Mechanism is the knowledge that data lie within some bounds. Specifically for the Elliptical Gaussian Mechanism data is required to lie within some hyperrectangle. It is formally described by equation (2) essentially saying that there is an upper and lower bound on each dimension. This requirement is needed to know the l_2-sensitivity $\Delta(f)$ as shown in equation (3). In this project I will change this assumption and instead look at the case where each dimension is normally distributed. This means that for each $j \in [d]$ we have that $X^{(j)} \sim \mathcal{N}(\mu_j, \sigma_j^2)$. An equivalent formulation is that a the data is multivariately distributed but with no correlation between dimensions. This means that $X \sim \mathcal{N}(\boldsymbol{\mu}, \Sigma)$, where Σ is a diagonal matrix with the variance of each dimension along its diagonal. It is quite apparent that determining a Δ_j is impossible in this setting as the

Gaussian distribution is continously defined on the range $(-\infty, \infty)$. Several recent papers has combatted this by doing something called *clipping* [3, 9]. Clipping is the process of limiting the norm of any one entry to be at most a chosen threshold C. This means that every vector is transformed as such

$$\widehat{x}_i := \min\left\{\frac{C}{\|x_i\|}, 1\right\} \cdot x_i$$

Clipping entries by a factor C thus means that $\Delta(f) = 2C$ as any one entry cannot have more impact on the summation than C. If the summation $f(X)$ is performed on a clipped dataset $\widehat{X}$ it is equivalent to defining the summation function $\widehat{f} : \mathbb{R}^{n \times d} \to \mathbb{R}^d$ as

$$\widehat{f}(X) = \sum_i^n \min\left\{\frac{C}{\|x_i\|}, 1\right\} \cdot x_i$$

Then by theorem 4 the gaussian mechanism with the function $\widehat{f}$ is (ε, δ)-DP with $\sigma_{\varepsilon, \delta} = \Delta\left(\widehat{f}\right)\sigma_{opt} = 2C\sigma_{opt}$. Though the mechanism is still (ε, δ)-DP it will now have a larger error when regarding the true sum $f(X) = \sum_i^n x_i$ as the actual answer. If the probability of clipping is set so low that we actually don't expect to clip any entries we can use $\widehat{f}$ as an approximation of f. Say there are n points in a dataset, I will thus set the probability of clipping to be less than $\frac{1}{n}$ and get that

$$\mathbb{E}\left[\left\|\left(\widehat{f}(X) + \eta\right) - f(X)\right\|^2\right] \approx \mathbb{E}\left[\left\|(f(X) + \eta) - f(X)\right\|^2\right] = \mathbb{E}\left[\|\eta\|^2\right]$$

The intuition in this context is the same as that for the Elliptical Gaussian Mechanism, that dimensions with high variance should have more noise added. Therefore I will use the same approach as that used by the Elliptical Gaussian Mechanism to achieve (ε, δ)-DP . They achieve (ε, δ)-DP by finding a transformation of points such that they all lie within the unit ball centered at the origin. There does not exist such a transformation which is linear, as some points will always lie outside the unit ball with at least a small probability. Instead I will introduce the constraint that the expected norm of vectors after the transformation is 1. I whish to find a scaling of each dimension b_j s.t.

Introduce σ_j before here

$$\mathbb{E}\left[\|x_i \odot \boldsymbol{b}\|\right] = 1 \iff \sum_{j \in [d]} (\sigma_j b_j)^2 = 1 \tag{6}$$

where $\boldsymbol{b} = (b_1, b_2, \ldots, b_d)$, and $\odot$ is the element-wise product. Then $\widehat{f}$ can be computed on this transformed dataset, where the probability of clipping is less than $\frac{1}{n}$. As Theorem 1 shows, (ε, δ)-DP is preserved under post processing, we can therefore add noise to each coordinate drawn from $\mathcal{N}(0, (2C\sigma_{opt})^2)$ in the transformed space to achieve (ε, δ)-DP . The transformation back to the original space is the done by multiplying each dimension with b_j^{-1}, and due to the linearity of transformation this is also done to the noise added. We end up with the following noise vector being added

$$\boldsymbol{\eta} = \left(b_1^{-1}\eta_1, b_2^{-1}\eta_2, \ldots, b_d^{-1}\eta_d\right)$$

in which all $\eta_j \sim \mathcal{N}(0, (2C\sigma_{opt})^2)$, and then by Lemma **??** we have $b_j^{-1}\eta_j \sim \mathcal{N}(0, (b_j^{-1} \cdot 2C\sigma_{opt})^2)$. As the error is approximately given by $\|\boldsymbol{\eta}\|$ we have due to Corollary 1.1 that the expected error is

$$\mathbb{E}\left[\|\boldsymbol{\eta}\|^2\right] = \sum_{j=1}^{d} \left(b_j^{-1} \cdot 2C\sigma_{opt}\right)^2 = (2C\sigma_{opt})^2 \sum_{j=1}^{d} b_j^{-2} \tag{7}$$

Now we are back to a very similiar problem to the one solved in [11]. Minimize the error described in equation (7) under the constraint given in equation (6). This leads to the following Lemma

Lemma 3.1. *Let $X \in \mathbb{R}^{n \times d}$ be a dataset where each dimension $X^{(j)}$ is independently gaussianly distributed, i.e. $X^{(j)} \sim \mathcal{N}(\mu_j, \sigma_j^2)$. Then the expected error of Algorithm **??** is minimized when*

$$b_j = \frac{1}{\sqrt{\sigma_j}\sqrt{\sum_{i=1}^{d} \sigma_i}}$$

Proof

Using lagrangian multipliers we find the local maxima or minia of the function subject to equality constraints. To do so we construct the lagrangian function $\mathcal{L} : \mathbb{R}^{d+1} \to \mathbb{R}$ from the optimization problem given by equation (7) and the constraint given by (6) for ease of notation we define $\sigma_{\varepsilon,\delta} := 2C\sigma_{opt}$.

$$\mathcal{L}(\boldsymbol{b}, \lambda) = \sigma_{\varepsilon,\delta}^2 \sum_{j \in [d]} b_j^{-2} + \lambda \left(\sum_{j=1}^{d} (\sigma_j b_j)^2 - 1\right)$$

We then find the stationary points of it, by setting the derivative of it to $\boldsymbol{0}$. The derivative with respect to b_j is

$$\frac{\partial \mathcal{L}}{\partial b_j}(\boldsymbol{b}, \lambda) = \frac{\partial}{\partial b_j}\left(\sigma_{\varepsilon,\delta}^2 \cdot b_j^{-2} + \lambda(\sigma_j b_j)^2\right) = -2\sigma_{\varepsilon,\delta}^2 b_j^{-3} + 2\lambda \sigma_j^2 b_j$$

I then solve $\frac{\partial \mathcal{L}}{\partial b_j} = 0$ for b_j

$$-2\sigma_{\varepsilon,\delta}^2 b_j^{-3} + 2\lambda \sigma_j^2 b_j = 0 \iff \lambda \sigma_j^2 b_j = \sigma_{\varepsilon,\delta}^2 b_i^{-3} \tag{8}$$

$$\iff b_j^4 = \frac{\sigma_{\varepsilon,\delta}^2}{\lambda \sigma_j^2} \iff b_j = \frac{\sqrt{\sigma_{\varepsilon,\delta}}}{\lambda^{\frac{1}{4}}\sqrt{\sigma_j}} \tag{9}$$

I now have the last partial derivative $\frac{\partial \mathcal{L}}{\partial \lambda} = 0$ which I solve for λ using the previous expression for b_j.

$$\frac{\partial \mathcal{L}}{\partial \lambda} = \sum_{j=1}^{d} (\sigma_j b_j)^2 - 1$$

$$\sum_{j=1}^{d} (\sigma_j b_j)^2 - 1 = 0 \iff \sum_{j=1}^{d} \sigma_j^2 \left(\frac{\sigma_{\varepsilon,\delta}}{\sqrt{\lambda}\sigma_j} \right) = 1 \iff$$

$$\frac{\sigma_{\varepsilon,\delta}}{\sqrt{\lambda}} \sum_{j=1}^{d} \frac{\sigma_j^2}{\sigma_j} = 1 \iff \sigma_{\varepsilon,\delta} \sum_{j=1}^{d} \sigma_j = \sqrt{\lambda}$$

Inserting back into equation 9

$$b_j = \frac{\sqrt{\sigma_{\varepsilon,\delta}}}{\lambda^{\frac{1}{4}}\sqrt{\sigma_j}} = \frac{\sqrt{\sigma_{\varepsilon,\delta}}}{\sqrt{\sigma_{\varepsilon,\delta} \sum_{i=1}^{d} \sigma_i}\sqrt{\sigma_j}} = \frac{1}{\sqrt{\sigma_j}\sqrt{\sum_{i=1}^{d} \sigma_i}}$$

Giving os the value of b_j which minimizes the expected error. ∎

3.1   Gaussian data

Let $X^{(j)} \sim \mathcal{N}(0, \sigma_j^2)$ As the expected l_2-norm of x_i is given by

$$\mathbb{E}\left[\|x_i\|^2\right] = \sum_{j=1}^{d} \sigma_j^2$$

To achieve an expected norm of 1 I will scale each dimension by a factor $\frac{1}{b_j}$ which achieves this. If

$$\hat{x}_i = \left(\frac{x_i^{(0)}}{b_0}, \frac{x_i^{(1)}}{b_1}, \dots, \frac{x_i^{(d)}}{b_d} \right)$$

This means that $X^{(j)} \sim \mathcal{N}(0, \frac{\sigma_j^2}{b_j^2})$ and the expected norm is given by

$$\mathbb{E}\left[\|x_i\|^2\right] = \sum_{j=1}^{d} \frac{\sigma_j^2}{b_j^2}$$

and I can introduce that constraint that the expected norm after the transformation should be 1. In such a case when noise is added after the transformation $\hat{X} + \eta$ where $\eta \sim N(0, t^2)$ and achieves (ε, δ)-DP in this space then then due to linearity of transformation the noise introduced in the original space is then given by then the error is

I desire a transformation of $x_i^{(j)}$ such that the expected norm is 1. Thus I must scale each dimension by $\frac{1}{b_j}$, and have that

$$\mathbb{E}\left[\|x_i\|^2\right] = 1$$

Minimize $\|\hat{\eta}\|$ under the constraint that $\mathbb{E}\left[\|x_i\|^2\right] = 1$

Lemma 3.2. *Let $X \sim \mathcal{N}(0, \sigma^2)$, and Φ denote the cumulative density function of $\mathcal{N}(0,1)$, then the cumulative density function of X^2 is given by*

$$F_{X^2}(x) = \Pr\left[X^2 \leq x\right] = 2\Phi\left(\frac{\sqrt{x}}{\sigma}\right) - 1$$

Proof

$$\Pr\left[X^2 \leq x\right] = \Pr\left[|X| \leq \sqrt{x}\right] = 2\Pr\left[0 \leq X \leq \sqrt{x}\right]$$
$$= 2\left(\Pr\left[X \leq \sqrt{x}\right] - \Pr\left[X \leq 0\right]\right) = 2\left(\Pr\left[X \leq \sqrt{x}\right] - \frac{1}{2}\right)$$
$$= 2\Phi\left(\frac{\sqrt{x}}{\sigma}\right) - 1$$

∎

Corollary 3.1. *From lemma 3.2 we can give following bound for $X \sim \mathcal{N}(0, \sigma^2)$.*

$$\Pr\left[X^2 > (4\sigma)^2\right] < 10^{-4}$$

Lemma 3.3. *Let $X_1, X_2, \ldots, X_d$ be d independent random gaussian variables where for $1 \leq j \leq d$ we have that $X_j \sim \mathcal{N}(0, \sigma_j^2)$. Then the probability for the sum of variables squared is bounded by*

$$\Pr\left[\sum_{j\in[d]} X_j^2 \geq t\sqrt{8\sum_{j\in[d]}\sigma_j^4} + \sum_{j\in[d]}\sigma_j^2\right] < e^{-t^2}$$

for

$$0 \leq t \leq \frac{1}{4 \cdot \max_{j\in[d]}\sigma_j^2}\sqrt{2\sum_{j\in[d]}\sigma_j^4}$$

Proof
At first we define the random variable $Y_j = X_j^2 - \mathbb{E}\left[X_j^2\right]$ using the j'th gaussian random variable. As $\mathbb{E}\left[Y_j\right] = \mathbb{E}\left[X_j^2 - \mathbb{E}\left[X_j^2\right]\right] = \mathbb{E}\left[X_j^2\right] - \mathbb{E}\left[X_j^2\right] = 0$ we have that Y_j is zero centered. We are thus interested in giving bounds on $\Pr\left[\sum_{j\in[d]} Y_j\right]$. We can use Bernsteins inequality, if the following constraint holds for all $k \in \mathbb{N}$ with $k \geq 2$ and for all $j \in [d]$, and for some $L \in \mathbb{R}$

$$\mathbb{E}\left[|Y_j^k|\right] \leq \frac{1}{2}\mathbb{E}\left[Y_j^2\right]L^{k-2}k! \tag{10}$$

Initially we have that $\mathbb{E}\left[|Y_j^k|\right] = \mathbb{E}\left[\left|\left(X_j^2 - \mathbb{E}\left[X_j^2\right]\right)^k\right|\right]$ and since $X_j^2 \geq 0$ and therefore also $\mathbb{E}\left[X_j^2\right] \geq 0$ we can therefore bound it by

$$\mathbb{E}\left[\left|\left(X_j^2 - \mathbb{E}\left[X_j^2\right]\right)^k\right|\right] \leq \mathbb{E}\left[|X_j^{2k}|\right] = \mathbb{E}\left[\left|\left(\sigma_j^2 Z\right)^k\right|\right] = \sigma_j^{2k} \cdot \mathbb{E}\left[|Z^k|\right]$$

9

Where $Z \sim \chi_1^2$ is a chi square with 1 degree of freedom. The moment generating function of Z is $\mathbb{E}\left[Z^m\right] = 1 \cdot 3 \cdot 5 \cdot \ldots \cdot (2m-1)$. Using this we can get the following bound

$$\sigma_j^{2k} \cdot \mathbb{E}\left[|Z^k|\right] = \sigma_j^{2k} \cdot \prod_{c=1}^{k}(2c-1) = \sigma_j^{2k} \cdot \prod_{c=2}^{k}(2c-1)$$

$$\leq \sigma_j^{2k} \cdot \prod_{c=2}^{k} 2c = \sigma_j^{2k} \cdot 2^{k-1} \cdot \frac{k!}{2} = \sigma_j^{2k} \cdot 2^{k-2} \cdot k!$$

Concluding that $\mathbb{E}\left[|Y_j^k|\right] \leq \sigma_j^{2k} \cdot 2^{k-2} \cdot k!$.
Secondly I calculate $\mathbb{E}\left[Y_j^2\right]$ exactly to be used in equation 10

$$\mathbb{E}\left[Y_j^2\right] = \mathbb{E}\left[\left(X_j^2 - \mathbb{E}\left[X_j^2\right]\right)^2\right] = \mathbb{E}\left[\left(X_j^2 - \mathbb{E}\left[X_j^2\right]\right)^2\right] = \tag{11}$$

$$\mathbb{E}\left[X_j^4 + \mathbb{E}\left[X_j^2\right]^2 - 2X_j^2\mathbb{E}\left[X_j^2\right]\right] = \mathbb{E}\left[X_j^4\right] + \mathbb{E}\left[X_j^2\right]^2 - 2\mathbb{E}\left[X_j^2\right]^2 = \tag{12}$$

$$\mathbb{E}\left[X_j^4\right] - \mathbb{E}\left[X_j^2\right]^2 = \sigma_j^4\mathbb{E}\left[Z^2\right] - \left(\sigma_j^2\mathbb{E}\left[Z\right]\right)^2 = \tag{13}$$

$$3\sigma_j^4 - \sigma_j^4 = 2\sigma_j^4 \tag{14}$$

Eqaution 10 is therefore rewritten to

$$\sigma_j^{2k} \cdot 2^{k-2} \cdot k! \leq \frac{1}{2} \cdot 2\sigma_j^4 L^{k-2} k! \iff$$

$$\sigma_j^{2k-4} \cdot 2^{k-2} \cdot \leq L^{k-2} \iff$$

$$\sigma_j^2 \cdot 2 \leq L$$

As this has to hold for all $j \in [d]$, we can simply choose $L := 2 \cdot \max_{j \in [d]} \sigma_j^2$ which satisfies the constraint.

We can then use Bernstein's inequality to bound the following

$$\Pr\left[\sum_{j \in [d]} Y_j \geq 2t\sqrt{\sum_{j \in [d]} \mathbb{E}\left[Y_j^2\right]}\right] < e^{-t^2} \tag{15}$$

Which in in this case can be rewritten to get a tail bound on $\sum_{j \in [d]} X_j^2$, by reusing results from equations (11)-(14)

$$\Pr\left[\sum_{j \in [d]} Y_j \geq 2t\sqrt{\sum_{j \in [d]} \mathbb{E}\left[Y_j^2\right]}\right] = \Pr\left[\sum_{j \in [d]}\left(X_j^2 - \mathbb{E}\left[X_j^2\right]\right) \geq 2t\sqrt{\sum_{j \in [d]} 2\sigma_j^4}\right]$$

$$= \Pr\left[\sum_{j \in [d]} X_j^2 \geq 2t\sqrt{\sum_{j \in [d]} 2\sigma_j^4} + \sum_{j \in [d]} \mathbb{E}\left[X_j^2\right]\right] = \Pr\left[\sum_{j \in [d]} X_j^2 \geq 2t\sqrt{\sum_{j \in [d]} 2\sigma_j^4} + \sum_{j \in [d]} \sigma_j^2\right]$$

$$= \Pr\left[\sum_{j \in [d]} X_j^2 \geq t\sqrt{8\sum_{j \in [d]} \sigma_j^4} + \sum_{j \in [d]} \sigma_j^2\right]$$

Finally giving us that

$$\Pr\left[\sum_{j\in[d]} X_j^2 \geq t\sqrt{8\sum_{j\in[d]}\sigma_j^4} + \sum_{j\in[d]}\sigma_j^2\right] < e^{-t^2}$$

where t must lie within

$$0 \leq t \leq \frac{1}{2L}\sqrt{\sum_{j\in[d]}\mathbb{E}\left[Y_j^2\right]} = \frac{1}{4\cdot\max_{j\in[d]}\sigma_j^2}\sqrt{2\sum_{j\in[d]}\sigma_j^4}$$

$\blacksquare$

Combining lemma 3.3 with theorem 3.1 we have can conclude the following:

Explain why α is not squared (it is the sum of squares and Δ is limited by $\|X\|$)

Theorem 6. *The expected error of algorithm 3 is given by*

$$\mathbb{E}\left[\|\eta\|^2\right] = \sigma_{\varepsilon,\delta}^2 \cdot \alpha \cdot \left(\sum_{i=1}^d \sigma_i\right)^2$$

where α is the smallest value statisfying $\Pr\left[\|X\|^2 > \alpha\right] < n^{-1}$.
When $\ln(n) \leq \frac{\sum_{j\in[d]}\sigma_j^2}{8\max_{j\in[d]}\sigma_j^2}$ the error can thus be upper bounded by

$$\mathbb{E}\left[\|\eta\|^2\right] \leq \sigma_{\varepsilon,\delta}^2 \cdot \sum_{j\in[d]}\sigma_j\left(\sqrt{8\ln(n)\sum_{i\in[d]}\sigma_i^2} + \sum_{j\in[d]}\sigma_j\right)$$

Proof
By equation 7 the error is given by

$$\mathbb{E}\left[\|\eta\|^2\right] = \sigma_{\varepsilon,\delta}^2 \cdot \Delta(\hat{f})^2 \cdot \left(\sum_{i=1}^d \sigma_i\right)^2 \tag{16}$$

When clipping is performed to remove less than n^{-1} we have the following

is this legal?

$$\Pr\left[\|X\| \geq \Delta(f)\right] = \Pr\left[\|X\|^2 \geq \Delta(f)^2\right] = \Pr\left[\sum_{j\in[d]} X_j^2 \geq \Delta(f)^2\right]$$

Which means I can give an upper bound on $\Delta(f)^2$ by using lemma 3.3, and inserting that $\sigma_j = \sigma_j \cdot b_i$ where b_i is given by theorem 3.1. I whish the clipping probability to be

less than n^{-1}, which implies $t = \sqrt{\ln(n)}$. Combining these results I get that

$$\Delta(f)^2 \leq \sqrt{8\ln(n) \sum_{j\in[d]} \left(\sigma_j^4\right) + \sum_{j\in[d]} \sigma_j^2}$$

$$= \sqrt{8\ln(n) \sum_{j\in[d]} \left(\frac{\sigma_j}{\sum_{i\in[d]} \sigma_i}\right)^2 + \sum_{j\in[d]} \frac{\sigma_j}{\sum_{i\in[d]} \sigma_i}}$$

$$= \sqrt{8\ln(n) \sum_{j\in[d]} \sigma_j^2 \cdot \frac{1}{\sum_{i\in[d]} \sigma_i} + 1}$$

Inserting this back into equation 16 we conclude

$$\mathbb{E}\left[\|\eta\|^2\right] = \sigma_{\varepsilon,\delta}^2 \cdot \Delta(\hat{f})^2 \cdot \left(\sum_{i=1}^d \sigma_i\right)^2$$

$$\leq \sigma_{\varepsilon,\delta}^2 \cdot \left(\sqrt{8\ln(n) \sum_{j\in[d]} \sigma_j^2 \cdot \frac{1}{\sum_{i\in[d]} \sigma_i} + 1}\right) \cdot \left(\sum_{i=1}^d \sigma_i\right)^2$$

$$= \sigma_{\varepsilon,\delta}^2 \cdot \sum_{i=1}^d \sigma_i \cdot \left(\sqrt{8\ln(n) \sum_{j\in[d]} \sigma_j^2} + \sum_{i=1}^d \sigma_i\right)$$

And the constraint on $t = \sqrt{\ln(n)}$ from lemma 3.3 is

$$0 \leq t \leq \frac{1}{4\max_{j\in[d]} \sigma_j^2} \sqrt{2 \sum_{j\in[d]} \sigma_j^4} \iff$$

$$\sqrt{\ln(n)} \leq \sqrt{\left(\frac{\sum_{j\in[d]} \sigma_j}{\max_{j\in[d]} \sigma_j}\right)^2 \cdot \frac{1}{8} \cdot \sum_{j\in[d]} \left(\frac{\sigma_j}{\sum_{i\in[d]} \sigma_i}\right)^2} \iff$$

$$\ln(n) \leq \frac{1}{8} \frac{1}{\max_{j\in[d]} \sigma_j^2} \cdot \sum_{j\in[d]} \sigma_j^2 \iff$$

$$\ln(n) \leq \frac{\sum_{j\in[d]} \sigma_j^2}{8\max_{j\in[d]} \sigma_j^2}$$

∎

Bernsteins inequality then states that

$$\Pr\left[\sum_{j\in[d]} X_j \geq 2t\sqrt{\sum_{j\in[d]} \mathbb{E}\left[X_i^2\right]}\right] < e^{-t^2}$$

for all

$$0 \leq t \leq \frac{1}{2L}\sqrt{\sum}$$

I have that $\bar{X}$ is a standard gaussian variable.

Let $Y_j = X_j^2$ I am interested in $\sum_{j\in[d]} Y_j$, then it has

$$\mathbb{E}\left[|Y_j^k|\right] = \mathbb{E}\left[|X_j^{2k}|\right] = \mathbb{E}\left[|X_j|^{2k}\right] = \mathbb{E}\left[|\sigma_j\bar{X}|^{2k}\right] =$$

$$|\sigma_j|^{2k}\mathbb{E}\left[|\bar{X}|^{2k}\right] = \sigma_j^{2k}\prod_{c=1}^{k} 2c = \sigma_j^{2k} \cdot 2^k \cdot k!$$

Which means $\mathbb{E}\left[Y_j^2\right] = 8\sigma_j^4$. To check the constraint (find value for L)

$$\sigma_j^{2k} \cdot 2^k \cdot k! \leq \frac{1}{2} \cdot 8\sigma_j^4 \cdot L^{k-2}k! \iff$$

$$\sigma_j^{2k} \cdot 2^k \leq 4\sigma_j^4 \cdot L^{k-2} \iff$$

$$\sigma_j^{2(k-2)} \cdot 2^k \leq 4\mathrm{L}^{k-2}$$

We have that $\sigma_j = \sigma_j b_j = \frac{\sqrt{\sigma_j}}{\sqrt{\sum_{i\in[d]}\sigma_i}} = \sqrt{\frac{\sigma_j}{\sum_{i\in[d]}\sigma_i}}$. Therefore

$$\sigma_j^{2(k-2)} \cdot 2^k = \left(\sqrt{\frac{\sigma_j}{\sum_{i\in[d]}\sigma_i}}\right)^{2(k-2)} \cdot 2^k =$$

$$\left(\frac{\sigma_j}{\sum_{i\in[d]}\sigma_i}\right)^{k-2} \cdot 2^k \leq 2^k$$

This therefore implies

$$2^k \leq 4L^{k-2} \iff 2^{k-2} \leq L^{k-2} \iff 2 \leq L$$

If it is desired that less than 10^{-p} points are removed then $t = \sqrt{p\ln(10)}$ as long as

$$\sqrt{p\ln(10)} \leq \frac{1}{2L}\sqrt{\sum_{j\in[d]} \mathbb{E}\left[X_j^2\right]} = \sqrt{4} \cdot \frac{\sqrt{\sum_{i\in[d]}\sigma_i^2}}{\sum_{i\in[d]}\sigma_i}$$

An alternative Again minimize $\|\hat{\eta}\|$, but instead the constraint comes from Chebyshev's inequality, I always have that

$$\Pr\left[|\|x_i\|^2 - \mathbb{E}\left[\|x_i\|^2\right]| \geq k \cdot \sqrt{\mathrm{Var}\left[\|x_i\|^2\right]}\right] \leq \frac{1}{k^2}$$

Therefore I can set $\frac{1}{k^2} = 0.05$, and find a transformation where I decide how many standard deviations I must be away from the mean to have norm greater than 1, i.e.

$$\mathbb{E}\left[\|x_i\|^2\right] + k \cdot \sqrt{\mathrm{Var}\left[\|x_i\|^2\right]} = 1 \implies k = \frac{1 - \mathbb{E}\left[\|x_i\|^2\right]}{\sqrt{\mathrm{Var}\left[\|x_i\|^2\right]}}$$

I then have that

$$\frac{1}{k^2} = \frac{1}{\left(\frac{1 - \mathbb{E}[\|x_i\|^2]}{\sqrt{\mathrm{Var}[\|x_i\|^2]}}\right)^2} = \left(\frac{\sqrt{\mathrm{Var}\left[\|x_i\|^2\right]}}{1 - \mathbb{E}\left[\|x_i\|^2\right]}\right)^2 = \frac{\mathrm{Var}\left[\|x_i\|^2\right]}{(1 - \mathbb{E}\left[\|x_i\|^2\right])^2}$$

I already know that

$$\mathbb{E}\left[\|x_i\|^2\right] = \sum_{j=1}^{d} \frac{\sigma_j^2}{b_j^2}$$

$$\mathrm{Var}\left[\|x_i\|^2\right] = 2\sum_{j=1}^{d} \frac{\sigma_j^4}{b_j^4}$$

I therefore have the constraint

$$\frac{\mathrm{Var}\left[\|x_i\|^2\right]}{(1 - \mathbb{E}\left[\|x_i\|^2\right])^2} = \frac{2\sum_{j=1}^{d} \frac{\sigma_j^4}{b_j^4}}{\left(1 - \sum_{j=1}^{d} \frac{\sigma_j^2}{b_j^2}\right)^2} = 0.05$$

Be aware that this could find cases where expected norm is greater than 1 and the constraint then says that they are never less than 1 in norm. Another restriction could be that expected norm is ¡ 1.

Chebyshev's inequality could also be used when expected norm should be 1 and then put a bound on number of std away one must be to have less than 0.05 fraction of data removed.

Extra

Determining α using Bernsteins inequality

$$\Pr\left[\|x_i\|^2 - \mathbb{E}\left[\|x_i\|^2\right] > t\right] < 2 \cdot \exp\left(-\frac{t^2/2}{\text{Var}\left[\|x_i\|^2\right] + C \cdot t/3}\right)$$

In the case where $\mathbb{E}\left[\|x_i\|^2\right] = 1$ and b_i is optimized in this case, we have that

$$P(\|x_i\|^2 > 1 + t) < 2\exp\left(-\frac{t^2/2}{\text{Var}\left[\|x_i\|^2\right] + C \cdot t/3}\right)$$

$$\text{Var}\left[\|x_i\|^2\right] = 2\frac{\sum_i^d \sigma_i^2}{\left(\sum_i^d \sigma_i\right)^2}$$

$$C = \max_{i \in [d]}(\sigma_i) \cdot \frac{16}{\sum_i^d \sigma_i}$$

C is decided such that less than 0.0001 fraction of the data is outside this bound in each dimension. i.e.

$$\forall j \in [d]: \ \Pr\left[X^{(j)} > C\right] < 0.0001$$

Solving for t

$$2 \cdot \exp\left(-\frac{t^2/2}{\text{Var}\left[\|x_i\|^2\right] + C \cdot t/3}\right) = 0.0001 \iff \ln(\frac{0.0001}{2}) = -\frac{t^2/2}{\text{Var}\left[\|x_i\|^2\right] + C \cdot t/3} \implies$$

$$t = -\frac{C\ln(\frac{0.0001}{2}) - \sqrt{\ln(\frac{0.0001}{2})(-18 \cdot \text{Var}\left[\|x_i\|^2\right] + C^2\ln(\frac{0.0001}{2}))}}{3}$$

Then we have that $\alpha = 1 + t$.

> Determine α for the bound using `https://en.wikipedia.org/wiki/Concentration_inequality`, or `https://web.stanford.edu/class/cs229t/2017/Lectures/concentration-slides.pdf`

References

[1] Balle, B., and Wang, Y. Improving the gaussian mechanism for differential privacy: Analytical calibration and optimal denoising. *CoRR abs/1805.06530* (2018).

[2] Bates, D. *Quadratic Forms of Random Variables.* University of Wisconsin-Madison: STAT 849 lectures, 2010.

[3] Biswas, S., Dong, Y., Kamath, G., and Ullman, J. Coinpress: Practical private mean and covariance estimation. *Advances in Neural Information Processing Systems 33* (2020), 14475–14485.

[4] Chen, T., and Lumley, T. Numerical evaluation of methods approximating the distribution of a large quadratic form in normal variables. *Computational Statistics & Data Analysis 139* (2019), 75–81.

[5] DWORK, C., MCSHERRY, F., NISSIM, K., AND SMITH, A. Calibrating noise to sensitivity in private data analysis. *Journal of Privacy and Confidentiality 7*, 3 (2016), 17–51.

[6] DWORK, C., ROTH, A., ET AL. The algorithmic foundations of differential privacy. *Foundations and Trends® in Theoretical Computer Science 9*, 3–4 (2014), 211–407.

[7] HALLAM, A. *Some Theorems on Quadratic Forms and Normal Variables.* Iowa State University: Econ 671 lectures, 2004.

[8] HANSON, D. L., AND WRIGHT, F. T. A bound on tail probabilities for quadratic forms in independent random variables. *The Annals of Mathematical Statistics 42*, 3 (1971), 1079–1083.

[9] HUANG, Z., LIANG, Y., AND YI, K. Instance-optimal mean estimation under differential privacy.

[10] MATHAI, A., AND PROVOST, S. Quadratic forms in random variables.

[11] PAGH, R., AND LEBEDA, C. Private vector aggregation when coordinates have different sensitivity.